

List of Commands:

You have been provided a zipped file that contains some folder/files etc for doing below activities. Download the file and unzip it (unzip filename ; replace filename with the name of whatever you downloaded). It should create a folder called student-files. Enter that folder. Most commands below assume you are inside the student-files folder.

Grep

1. Grep: grep searches for PATTERNS in each FILE and prints each line that matches that pattern. Note line is not decided by full-stops in the para (like we humans do) but by a newline (\n). "-i" ignores the case. The first command will not show a sentence starting with "For", whereas the second command will. The third command matches the exact pattern made of two words.
 - a. `grep "for" bigfile`
 - b. `grep -i "for" bigfile`
 - c. `grep -i "for a" bigfile`
2. Regex Metacharacters (^ beginning of line; \$ end of line; . match any single character)
 - a. `grep -i "^for" bigfile`
 - b. `grep -i "cried.$" bigfile`
 - c. `grep -i "h.s" bigfile`
 - d. `grep -i "t.t" bigfile`
3. More metacharacters: Suppose you want to use "or" operator i.e. "|". The first command will fail, since it tries to match the same exact pattern. The second works, since by using \, you are escaping the | i.e you will interpret it as the "or" operator. But this tends to be not clear when reading regular expressions. The best way around is the use of the third command, use of -E (has to be capital) in which case you don't need to escape special characters. Use the man page of grep to know more.
 - a. `grep "cried|could" bigfile`
 - b. `grep "cried\|could" bigfile`
 - c. `grep -E "cried|could" bigfile`
4. Quantifiers: * zero or more times; ? zero or one time; + one or more times; {n} exactly n times; {n,} at least n times; {,m} at most m times; {n,m} from n to m times
 - a. `grep -E "to*" bigfile`
 - b. `grep -E "to?" bigfile`
 - c. `grep -E "to+" bigfile`
 - d. `grep -E "to{1}" bigfile`
 - e. `grep -E "to{1,}" bigfile`
5. Groups and Ranges: () group patterns together; { } match a particular number of occurrences (saw this in 4.d and 4.e above); [] match any character from a range of characters. In the third command, the interpretation of "^" is not beginning of a line, rather since it appears within [], it means any letter except i. The fourth command, many

any letter in the alphabet range from A to Z (notice capitals, it will only match capital letters).

- a. `grep -E "(fear)?less" bigfile`
 - b. `grep -i "h[ia]s" bigfile`
 - c. `grep 'h[^i]s' bigfile`
 - d. `grep "[A-Z]" bigfile`
6. Special characters. `\s` indicates white space. So, the first command will search for a three letter word that starts with h and ends in d; which has a space in front and at the end. `\b` matches any character that is not a letter or number without including itself in the match.
- a. `grep -E "\sh.d\s" bigfile`
 - b. `grep -E "\bh.d\b" bigfile`
7. Grep is not restricted to just file search
- a. `ls -l | grep file`
8. More options of grep: `-v` negation, `-r` recursive search, `-w` whole word, `-n` show line number, `-c` count (count of lines, not number of occurrences). In the first command, every line that has the word "Wolf" is not printed! In the second command, only those lines that have "Wolf" are printed. In third command, we are using `-r` to search across all files in that folder (specified via `*`) i.e grep can work across files also, need not limit to one file.
- a. `grep -v "Wolf" bigfile`
 - b. `grep "Wolf" bigfile`
 - c. `grep -r "Wolf" *`
 - d. `grep -w "his" bigfile`
 - e. `grep "his" bigfile`
 - f. `grep -n "Wolf" bigfile`
 - g. `grep -c "Wolf" bigfile`

find

1. Locate files/directories with known pattern. Ensure you are in the parent of student-files folder. The first command, `-name` tells to match the name of the given file (i.e. files ending in jpg) in the folder student-files. The second command is asking it to check only for files (`-type f`) ending in ".c" and be case insensitive in file names (`-iname`) in the current directory.
 - a. `find student-files/ -name "*.jpg"`
 - b. `find . -type f -iname "*.c"`
2. The first command lists all directories in the current directory (`-type d`)
 - a. `find . -type d`
3. The first command finds files of size more than 60kbytes in current directory (use `"ls -lRl ."` to check the file sizes to verify if the command worked fine). The second command shows files with permission set to 644 in the current folder (permissions will be covered

shortly). The last command, finds files starting with f in dir1 folder and then deletes the files within (be very careful when you execute this, since it will remove files)

- a. `find . -type f -size +60k; ls -RI .`
- b. `find . -perm 644`
- c. `find dir1/ -name "f*" -delete`

wc

1. Ensure you are inside student-files folder. wc prints word (-w), character (-m), line (-l), byte-count (-c) in a file. wc can accept zero (reads from standard input) or more inputs. Default output: lines, words, characters. The first command prints the lines, words and characters (in that order) to the output of file bigfile. The second command prints only number of lines. Try the remaining commands also, they are straightforward.
 - a. `wc bigfile`
 - b. `wc -l bigfile`
 - c. `wc -w bigfile`
 - d. `wc -m bigfile`
 - e. `wc -c bigfile`
2. We can use wc with multiple files (below we are using it across two files)
 - a. `wc /proc/cpuinfo /proc/meminfo`

cut:

1. cut: helps cut parts of a line by delimiter, byte position, and character ("-f" specifies field(s), "-b" specifies bytes(s), "-d" specify a delimiter, --complement complement the selection and --output-delimiter specify a different output delimiter string)
2. Below are some examples involving students.csv file. If you do `cat students.csv`, you will see various fields separated by comma (which is the delimiter). The first command specifies what delimiter to use (, in this case) and to output fields 1 and 3 (serial no and name). The second command is using -4 which means fields 1 to 4.
 - a. `cut students.csv -d ',' -f 1,3`
 - b. `cut students.csv -d ',' -f -4`
3. You can change the output delimiter as well. First command now separates the field by space when outputting.
 - a. `cut students.csv -d ',' -f 1,3 --output-delimiter=" "`
4. When using complement, it outputs fields that are complement of 1,2,3 i.e 4,5,6,7,8
 - a. `cut students.csv -d ',' -f 1,2,3 --complement`
5. Here are examples of -b. Note I am using echo and redirecting its output to cut. You don't necessarily have to use a file
 - a. `echo "abcdefg" | cut -b 1,3,5`
 - b. `echo "abcdefg" | cut -b 1-4`

paste:

1. Helps merge lines of files horizontally. -d can be used to specify a delimiter when pasting. And -s pastes one file at a time in serial.
 - a. `paste fruits1 fruits2`
 - b. `paste -d '_' fruits1 fruits2`

- c. `paste -s fruits1 fruits2`

sort

1. sorts the records in a file. First command sorts in ascending order. And second command sorts descending (reverse)
 - a. `sort fruits1`
 - b. `sort -r fruits1`
2. When you have numbers, sort will still character wise without taking the numerical value under condition i.e. (13 will come before 9). Use `-n` to sort numerically. In third command, we are combining two options, `-n` and `-r` as `-nr`.
 - a. `sort list1`
 - b. `sort -n list1`
 - c. `sort -nr list1`
3. One can also sort by specifying a field separator via `-t` option and key via `-k` option. Below will sort the data in `students.csv` based on 8th field i.e. key (hostels). Note separator is needed to tell what the fields are.
 - a. `sort -t ',' -k8 students.csv`

zip

1. zip archives files with compression. The first command recursively zips all contents of `fun_dir` as well as a file `students.csv`. The second command will encrypt the zip via `-e` option, it will ask for a password. The
 - a. `zip -r example.zip fun_dir/ students.csv`
 - b. `zip -r example.zip fun_dir/ students.csv`
2. unzip unarchives. The `-l` option does not decompress but lists what is inside the zip including file sizes, modification dates etc. Useful to know before unzipping. The second command unzips the given zipped file into a specified directory, which is specified via the `-d` option. If you instead just typed “unzip example.zip”, it would unzip in the current directory itself (if similar folder/files are present, it will ask if it should overwrite).
 - a. `unzip -l example.zip`
 - b. `unzip example.zip -d dir1/`

ps and kill

1. ps displays a list of processes currently executing. When you type ps in a shell without any options, it will show only the processes associated with that shell. In the second command, we are using options, `aux`, “a”: display the processes of all users; “u”: user-oriented formats; “x”: list processes without a controlling terminal
 - a. `ps`
 - b. `ps aux`
2. ps command also allows you to sort the output. For example, to sort the output based on the memory usage, you would use:
 - a. `ps aux --sort=-%mem`

3. kill can terminate a process. It is typically used if you find some process not responding
To kill a process, we need its process id. Type below in sequence. The first command is starting a process that will just sleep for 60sec, & instructs it to run in background, so the terminal is freed up (you can type sleep 60 without & and see what happens). The second command shows the pid of sleep. Then use that pid to kill via the third command. Then type ps again to see that it is indeed killed.
 - a. sleep 60 & ; ps; kill pid-of-sleep; ps

Redirection >, >> and <

1. > redirects the output of a command into a file. >> appends instead of overwriting
 - a. ls -l > ls-out
 - b. ls >> ls-out; cat ls-out
 - c. ls > ls-out; cat ls-out
2. When a program prints both to the stdout and also has some error messages printed to stderr, you can use fd> to redirect output accordingly. commands.sh is a program (bash script, which we will cover later, for now you treat it just as a black box program; you have). The first command will direct the output to the out file, the error messages are printed on the screen. The second command directs the error messages to out file, while the output is print to stdout. The third command sends both to two different files. The fourth command sends both to same file. The fifth command sends the error messages to a special file called /dev/null that discards anything written to it (used to suppress error messages)
 - a. ./commands.sh 1> out
 - b. ./commands.sh 2> out
 - c. ./commands.sh 2> error 1> out
 - d. ./commands.sh > out 2>&1
 - e. ./commands.sh 2> /dev/null
3. You can use < to read from a file instead of stdin. If you just type cat, whatever you type at a terminal it will echo back. You can get out by typing ctrl+c. But if you use cat < smallfile it will that the input from that file (note "cat smallfile" will also work as such)
 - a. cat
 - b. cat < smallfile
4. You can use both > and < also
 - a. wc -l < smallfile > lines

Pipe |

1. Pipe takes output from one command and feeds it as input to another command. The first command feeds the output of ls to wc. The second command uses two pipes and essentially tells the number of (sub)directories in the current directory. The third command also uses two pipes. The cat combines two files and feeds it to sort, which then feeds it to a command uniq which removes duplicates (-i asks it to ignore case)
 - a. ls /etc | wc -l
 - b. ls -l | grep "^d" | wc -l
 - c. cat fruits1 fruits2 | sort | uniq -i

Command Substitution

1. Output of a command replaces the command itself. Usage `$(command)` or ``command``.
2. We will cover this via an example using `date` and `touch` commands. `date` is a command which tells the current date. With option `+%s`, it tells the seconds passed since January 1st, 1970 at 00:00:00 UTC, which is referred to as the Unix epoch (standardized reference point for measuring time across different systems). `touch` is a command which can be used to create empty files.
 - a. `date +%s`
 - b. `touch file.txt ; ls`
3. Suppose you want to uniquely name files, you could use the `date` command to timestamp the file as follows. The output of the `date` command is now part of the `touch` command. Both the below commands have the same result, though timestamp will be different since time of execution of commands is different.
 - a. `touch file-`date +%s`.txt ; ls`
 - b. `touch file-$(date +%s).txt ; ls`

Access Control

1. Permissions for a file or directory may be any or all of : `r` - read; `w` - write; `x` - execute. Access permissions are displayed using `ls -l`. Observe a few things, `commands.sh` has `rxw` permissions i.e. why it could be executed via `./commands.sh`. The file `oddball` has read permissions but no write permissions. `Cat` works but open it in `nano` and try to write, you will not be able to write!
 - a. `ls -l`
 - b. `cat oddball`
 - c. `nano oddball`
2. Let us change some permissions now. First command uses the numeric mode to remove all permissions for all users. Neither `cat` nor `nano` to write will work now. The second command uses symbolic mode to change permission of a user to read. Now the `cat` command should work but cannot edit via `nano` still. The third command uses numeric mode to assign both read and write permissions to the file. Now both read and write should work.
 - a. `chmod 000 oddball; cat oddball; nano oddball`
 - b. `chmod u=r oddball ; cat oddball; nano oddball`
 - c. `chmod 640 oddball; cat oddball; nano oddball`
3. Another example which removes read permissions from a group recursively from a given folder. Go through the files before and after and see the diff.
 - a. `ls -lR fun_dir > before ; chmod -R g-r fun_dir; ls -lR fun_dir > after`

1. apt helps install, update, remove packages on Ubuntu, Debian etc. Most of the apt commands need sudo privileges. The first command shows details of the package rolldice. This does not need sudo permissions. The second command installs the package, needs sudo permissions. The third command is calling the installed package. You need to press enter to see dice roll output and do ctrl+C to exit. The fourth command removes the package from the system. When you type rolldice again, it says no such file or directory, since it is removed.
 - a. apt show rolldice
 - b. sudo apt install rolldice
 - c. rolldice
 - d. sudo apt remove rolldice
 - e. rolldice